

Energy-Sensitive Vision-Based Autonomous Tracking and Landing of a UAV

Georgios Zamanakos, Adam Seewald, Henrik Skov Midtiby, and Ulrik Pagh Schultz
SDU UAS Center, Mærsk Mc-Kinney Møller Institute, University of Southern Denmark
Contact email: ups@mmmi.sdu.dk

Abstract

In this paper we present a robust, vision-based algorithm for autonomous tracking and landing on a moving platform in varying environmental conditions. We use a novel landing marker robust to occlusions to track the moving platform and the YOLOv3-tiny CNN to detect ground-based hazards in an agricultural use case. We perform all computations on-board using an *Nvidia Jetson Nano* and analyse the impact on the flight time by profiling the energy consumption of the landing marker detection algorithm and YOLOv3-tiny CNN. Experiments are conducted in Gazebo simulation using the *powprofiler* energy modeling tool to measure the energy cost as a function of Quality of Service (QoS). Our experiments test the energy efficiency and robustness of our system in various dynamic wind disturbances. We show that the landing marker detection algorithm can be run at the highest QoS with only a marginal energy overhead whereas adapting the QoS level of YOLOv3-tiny CNN results in a considerable power saving for the system as a whole. The power saving is significant for a system executing on a fixed-wing UAV but only marginal if executing on a standard multirotor UAV.

28

I. INTRODUCTION

29 Unmanned Aerial Vehicles (UAVs) are increasingly used for applications
 30 such as monitoring, surveillance, transportation of small payloads, and
 31 agricultural applications [1], [2]. One of the major constraints of such
 32 applications is their limited level of autonomy due to battery limitations.
 33 Extending the flying time of a UAV is normally done by having it land
 34 in order to replace or charge the battery before continuing the mission.
 35 Performing landings autonomously can however be challenging depending
 36 on the environment and whether the landing platform is stationary or
 37 mobile. Moreover, relying solely on the availability of a GPS signal for
 38 autonomous precision landing is not considered safe, since GPS signals
 39 can be temporarily lost or even tampered with. As an alternative, in this
 40 paper we investigate the use of a novel vision-based autonomous landing
 41 system, and evaluate its robustness towards environmental conditions such
 42 as visual disturbances and wind.

43 Extension of the flight time can be also achieved by using *energy-*
 44 *sensitive algorithms* that can reduce energy consumption by reducing the
 45 quality of service (QoS). With this approach, energy-costly computations
 46 such as computer vision are adapted by selecting the desired quality of
 47 service to match the available energy [3]. By combining energy-sensitive
 48 algorithms with autonomous landing capabilities, we aim to increase the
 49 total availability of the UAV to perform operations, by extending the flight
 50 time and using autonomous recharging when needed.

51 The main contribution of this paper concerns the experimental study
 52 of a robust vision-based algorithm for autonomous tracking and landing
 53 in varying environmental conditions. The algorithms are executed on an
 54 NVIDIA Jetson Nano companion computer controlling a simulated drone.
 55 The vision-based tracking and landing algorithms provide novel capabili-
 56 ties in terms of tolerance to visual disturbance and varying environmental
 57 conditions such as wind. Our experiments are based on an agricultural use
 58 case where a multirotor UAV performs visual identification of ground-

59 based hazards while tracking and landing on a moving platform.

60 II. STATE OF THE ART

61 Vision-based autonomous landing on a marker has been extensively
62 studied by many researchers. Key distinctions include whether the marker
63 is on a moving platform, the type of the marker, the algorithms used to
64 detect it, as well as the mounted sensors on-board of the UAV.

65 For stationary platforms, one of the first experiments with vision-based
66 autonomous landing was conducted by Saripalli et al. [4]. Here, a heli-
67 copter with a color camera facing vertically towards the ground would
68 land on an H-shape pattern (similar to ones found on a helipad) using a
69 hierarchical behavior-based control architecture. In physical tests a marker
70 of $122\text{ cm} \times 122\text{ cm}$ size was detected for a maximum altitude of 10 m.
71 A landing marker inspired by a QR code but consisting of three artificial
72 markers is demonstrated by Yuan et al. [5], and was shown to provide a
73 6-Degree Of Freedom (DOF) pose over an altitude range of 0-20 m. Our
74 work is however focused on the ability to land on moving platforms.

75 Saripalli et al. [6] also demonstrated the use of a Kalman Filter to
76 track a moving platform. However all the computations were performed
77 offline. Similarly, an ArUco marker was used as a landing marker by
78 Lee et al. [7] to detect a moving platform. The control of the UAV is
79 performed based on the error provided by the vision algorithm but all the
80 computations were performed off-board. Arrar, et al. [8] focus on extending
81 the detection range by using an AprilTag [9] as a landing marker. Again all
82 the computer vision algorithms were also executed off-board. Conversely,
83 a crucial aspect of our application is to perform all the computations
84 on-board, and to evaluate them according to their energy efficiency as
85 a function of QoS.

86 The design of the marker and choice of sensors can facilitate doing the
87 computations on-board. Chen et al. [10] utilized a marker consisting of
88 a circle and rectangles of different colors along with a LiDAR scanning

range finder for height estimation. The marker was detected by performing color segmentation on the incoming image frame. By fusing the height measurement from the LiDAR into the vision measurement, a relative pose of the UAV from the moving platform was obtained. A color segmentation approach was also implemented by Lee et al. [11]. A red rectangle was used as a landing marker while the detection was done by a vertically facing camera with a fish-eye lens. The setup accounted for a successful landing from an altitude of 70 m.. Both researchers have used an on-board companion computer to perform all the computations on the UAV. However we in this work do not consider a color segmentation approach as a safe option, since for a realistic (outdoor) case it would be difficult, if not impossible, to ensure that the landing marker will be the only object of a specific color in the scene.

The use of a hybrid camera system consisting of a fish-eye IR camera and a stereo camera was demonstrated by Yang et al. [12]. An ArUco marker was used to mark the moving platform and a convolutional neural network (CNN) YOLO v3 was trained specifically for marker detection. A similar approach concerning the detection of a landing marker was demonstrated by Nguyen et al. [13]. Here a specific landing marker was used and a specific CNN was used to detect it: successful detection of a $1\text{ m} \times 1\text{ m}$ marker size was demonstrated from a distance of 50 m. An AprilTag marker was used as a landing marker by Kyritsis et al. [14]. The identification of the AprilTag marker was performed through Graphics Processing Unit (GPU). In the above three cases, the researchers have utilized the companion's computer GPU to detect the landing marker. In the agricultural use case addressed in this paper, the GPU is however needed for a CNN to detect ground hazards, and since the GPU cannot normally run different algorithms simultaneously, the CPU should be used for detecting the landing marker. By doing so, a different QoS can be chosen for each algorithm.

To account for the energy modeling of computer vision algorithms, we

considered the work previously carried by Nardi et al. [15]. The authors present SLAMBench, a framework that investigates Simultaneous Localisation and Mapping (SLAM) algorithms configuration alternatives for energy efficiency. In our work we use `powprofiler`, a generic energy modeling tool [16]. This tool enables measuring the energy impact of different configurations of the ROS-based system implementing the agricultural use-case. In this paper we present extensions to `powprofiler` that facilitates the initial exploration of the energy usage of complex ROS-based systems.

Other approaches to energy modeling, such as the mission-based energy models studied by Sadrpour et al. [17], [18], focus mostly on ground-based autonomous vehicles instead of the UAVs. Morales et al. [19] extensively investigated the relation between motion and energy in a robot, but do not account for on the energy required for computation. Energy modeling of mobile robots as carried by Mei et al. [20]–[22] has provided the ground for the concept of modeling computation for energy-sensitive algorithm design. Indeed, the approach employed in this paper has evolved from an energy-efficient motion planning technique in [22], a design strategy that allows accounting for motion and computations separately in [21], to an energy-efficient deployment algorithm in [20].

The battery in our system is considered in the context of a drone being able to perform its mission while accounting for the eventuality of a battery shortage; to this end, we investigated the approach presented by Berenz et al. [23], where a battery management mission-based dynamic recharge approach is presented. A set of recharge stations are used, along with self-docking capable robots. Our approach similarly allows landing on a moving platform for recharging, which is in the context of this paper is considered in the proximity of the drone. The actual landing is handled by the proposed algorithm, and we also account for the energy required for executing this algorithm during landing.

Taking into account varying environmental conditions and unpredictable

151 movements of the platform to land on is relevant for the use of landing
 152 in outdoor, mobile scenarios. Regarding wind conditions, an AprilTag
 153 marker was used by Feng et al. [24] with a constant wind speed of 5 m/s
 154 as an external disturbance in a simulation environment. Nevertheless, a
 155 fluctuation in the wind’s magnitude and direction is likely to happen in
 156 realistic cases. Concerning estimation of the moving platform’s position
 157 and velocity, similar to our approach a Kalman Filter or Extended Kalman
 158 Filter (EKF) has been used for the estimation [8], [24], [25], whereas
 159 Yang et al. [12] constructed a velocity observer algorithm by calculating
 160 the actual moving distance of the moving platform over a period of time.

161 III. ENERGY-SENSITIVE MISSION DEPLOYMENT

162 A. Overall approach

163 The energy-sensitive design is a mission-oriented concept that adjusts
 164 the computations to the mission being performed while taking into account
 165 energy requirements, including energy consumed by actuation, computa-
 166 tion, and the presence of a limited power source. Specifically, in our agri-
 167 cultural use-case, the concept is employed to profile and eventually adapt
 168 the computationally heavy algorithms performing autonomous tracking,
 169 landing, and hazard detection. This adaptation enables energy-sensitivity,
 170 in the sense that QoS parameters can be modified to enable the mission
 171 to be completed at the highest possible QoS level that does not exceed
 172 the available energy budget. Tradeoffs between QoS parameters can be
 173 performed by an end-user, i.e., trading the robustness towards wind during
 174 landing for precision of hazard detection.

175 The energy-sensitive design using `powprofiler` relies on empirical
 176 experiments to measure the actual power consumption on the robot hard-
 177 ware [16]. (The `powprofiler` tool is part of the TeamPlay toolchain,
 178 which aims to make tradeoffs between energy and other non-functional
 179 properties accessible to the developer.) In this paper we focus on the initial
 180 profiling using of energy usage of the companion computer, which from

the point of view of energy consumption can be studied independently from the specific drone it is mounted in.

First, the developer specifies the maximum and minimum QoS level for each algorithm running on the system. During mission execution the levels are statically defined: automatic adaptation during different phases of the mission is currently being investigated and is considered future work. Then, the developer executes the system to empirically determine the power consumption. This can be done in two different ways:

- 1) Automatically using `powprofiler` to control the experiment execution [16]. For a ROS-based system, we assume that the algorithms are wrapped as ROS nodes, and we require the developer to specify the QoS parameters using a ROS configuration. We use a configuration file in a key-value pair format which is then interpreted by `powprofiler`, enabling `powprofiler` to iterate through all possible combinations and empirically sample the energy consumption of each combination of QoS parameters. Once all combinations have been iterated through, `powprofiler` automatically combines the energy consumption data into a complete model.
- 2) Semi-automatically using `powprofiler` to sample energy and combine the results of all experiments, but allowing the developer to control all aspects of the experiment execution. This approach is new and is described in more detail later in this section. Basically, a ROS node interfaces to `powprofiler` and is used by the developer to start/stop sampling in a given configuration. Once all experiments have completed, `powprofiler` is invoked by the developer to combine the energy consumption data into a complete model.

Regardless of the approach, `powprofiler` builds a single model mapping QoS to total system energy consumption. Coarse-grained sampling is employed to reduce the number of experiments, and missing values are automatically inferred from the others by the means of a multivariate linear interpolation.

212 In the context of this paper, sampling experiments are iterated in a
 213 simulated environment with different configurations. For example, the au-
 214 tonomous tracking allows changing the tracking algorithm QoS in terms
 215 of frequency, the landing algorithm in terms of frequency, and hazard
 216 detection QoS in terms of frequency.

217 *B. Semi-automatic energy profiling*

218 A ROS node has been developed for the purposes of the semi-automatic
 219 approach described in this paper. This node allows automatic generation
 220 of the basic energy models that map time to the instantaneous power con-
 221 sumption. To activate this functionality, the developer simply publishes on
 222 a ROS topic to start the model generation, with `powprofiler` accounting
 223 for the invocation of an asynchronous thread which collects data from the
 224 energy sensors. Similarly, the developer publishes on another ROS topic
 225 to stop the model generation, while `powprofiler` finalizes collecting
 226 data from sensors, builds the basic energy model, and stores it for later
 227 processing.

228 Once all the basic energy models for the desired QoS ranges have been
 229 collected, QoS ranges are specified in a configuration file: the developer
 230 defines what QoS configuration corresponds to which basic model (instan-
 231 taneous power consumption as a function of time). Running `powprofiler`
 232 using this configuration file as a parameter generates the complete model
 233 that maps QoS to energy consumption.

234 IV. VISION-BASED AUTONOMOUS TRACKING AND LANDING

235 The vision-based autonomous tracking and landing can be split into
 236 four main sub-problems: detection of the moving platform, navigation,
 237 guidance, and control of the UAV. From an energy-sensitive design ap-
 238 proach, our focus is on the computer vision algorithms used to detect the
 239 moving platform and the parameterization by a QoS influencing energy
 240 consumption and performance, as described in Section IV-A. Furthermore,
 241 the navigation block is designed to increase the robustness and overall




Fig. 1. The landing marker

performance of the system, as described in Section IV-B. Last, a model of dynamically changing wind disturbances is analysed and described in Section IV-C, allowing the system to be tested on a realistic use case.

A. Detection of the moving platform

To mark the moving platform a special pattern is constructed, consisting of an n-fold marker [26] along with three ArUco markers [27] with different ids. This pattern will be referred to as the *landing marker* and can be seen in Figure 1. The n-fold marker is primarily used to detect the moving platform from a high altitude, while the ArUco markers are used as extra landmarks in case the marker is partially visible in the image frame.

To evaluate the computer vision algorithm for detecting the landing marker, real images of 640×480 pixel size were captured with an Intel RealSense D435 camera. In the Gazebo simulation a color camera with the same distortion coefficients as the Intel camera is used to output a 640×480 pixel image at 10 frames per second (fps).

To extract the pixel coordinates of the tip of the n-fold marker, a kernel size of 13×13 pixels consisting of a real and imaginary part is created. For every pixel in the image, a convolution is performed with this kernel and the magnitude of the convolution is stored. The pixel with the highest magnitude is considered as a candidate tip of the n-fold marker. For that




Fig. 2. Detection of the landing marker under different occlusions. On the left an occlusion on the tip of the n-fold marker and on the right an occlusion on a sector of the n-fold marker.

262 candidate pixel only, an estimation of the orientation/phase of the marker
 263 is made and an overall normalized quality score between 0.0 and 1.0 is
 264 calculated. If the score is above a desired threshold value then the pixel is
 265 accepted as the tip of the n-fold marker. If an n-fold marker is detected,
 266 the result will be the pixel coordinates of the tip of the n-fold marker along
 267 with its orientation/phase.

268 Since a convolution is a computationally expensive process, an increase
 269 in the kernel size would also increase the computation time and therefore
 270 the energy consumption. However, a higher computation time and energy
 271 consumption is preferred over a non-detection of the n-fold marker. To
 272 balance between energy consumption and effective marker detection, an
 273 adaptive kernel selection function is created to ensure the selection of a
 274 proper kernel size based on a threshold quality score value. In Figure 2
 275 an Intel RealSense D435 Depth camera is used to capture the image and
 276 measure also the distance d from the n-fold marker. Based on a desired
 277 quality q , the proper kernel size k is selected. It is seen that an occlusion on
 278 the tip of the n-fold marker results in a significant increase in the selected
 279 kernel size.

280 To detect the ArUco markers the standard OpenCV library is used and
 281 if any ArUco markers are detected, their central pixel coordinates and pose
 282 are stored.

283 The next step is to convert these pixel coordinates into a real-world

relative position $[X, Y, Z]$ according to a local coordinate frame. The origin of the local coordinate frame $[0, 0, 0]$ is defined as the center of the landing marker and alignment according to the North, East, Down (*NED*) frame. The available sensor measurements and sensor fusion algorithms from the flight controller are used in this process. The PX4 flight controller outputs through mavlink messages the altitude and the attitude of the UAV (roll, pitch, yaw). For the Z component, the altitude of the UAV from the flight controller's EKF is used. To obtain the X, Y components, an algorithm is constructed to convert pixel coordinates into real world X, Y coordinates in meters:

- 1) The pose of the camera in UAV's *BODY* frame is calculated by utilizing the roll and pitch IMU data.
- 2) The normalized coordinates of the four image corners, according to the camera's horizontal, vertical field of view and the camera's pose from step 1, are calculated.
- 3) The coordinates of the four image corners (in meters), with respect to the UAV's *BODY* frame, are determined by using the normalised coordinates from step 2 along with the UAV's altitude. The result is a projection plane of the image corners on the ground.
- 4) The perspective homography matrix is calculated between the 2 planes, the image plane and the world plane from step 3.
- 5) The homography matrix from step 4 is used to convert the pixel coordinates of the tip of the n fold marker from the image plane, into coordinates (in meters) in the UAV's *BODY* frame.
- 6) The coordinates from step 5 are in respect to the UAV's *BODY* frame. To convert them into the UAV's *NED* frame, the yaw IMU data from the flight controller is used.
- 7) The coordinates from step 6, are converted from UAV's *NED* frame into the landing site's local coordinate frame.
- 8) For ArUco markers only, an offset vector in the x, y axis is added depending on the distance of each ArUco marker from the tip of the

315 n-fold marker. It is assumed that this vector is prior known.

316 The mean measurements from the detected n-fold and/or ArUco markers
 317 are used to determine the position and orientation of the landing marker.
 318 The result is an $[X, Y, Z]$ relative position of the UAV from the moving
 319 platform along with the yaw orientation of the landing marker.

320 *B. Navigation*

321 The purpose of the navigation block is to provide an accurate prediction
 322 for the state of the UAV at any given time. This is important because it
 323 will allow us to process images at different fps according to a desired
 324 QoS. Furthermore, the overall robustness of the system is increased in
 325 case the moving platform is not detected in every image frame. A velocity
 326 estimator for the moving platform will also be implemented as a part of
 327 the navigation block.

328 The variables of interest that describe the state of the UAV for this
 329 project are:

- 330 • The relative position of the UAV from the moving platform $[X, Y, Z]$
 331 calculated as described in Section IV-A.
- 332 • The attitude of the UAV [roll, pitch, yaw], obtained from the flight
 333 controller's IMUs.
- 334 • the velocity of the UAV $[v_x, v_y, v_z]$ in *NED* frame, obtained from the
 335 flight controller's EKF.
- 336 • the acceleration of the UAV $[a_x, a_y, a_z]$ in *NED* frame, obtained either
 337 from the flight controller's EKF or by differentiating the velocities.

338 The altitude, attitude, velocity and acceleration variables of the UAV's
 339 state are obtained from the flight controller's onboard sensors and already
 340 implemented sensor fusion algorithms (EKF). To fuse those state variables
 341 with the obtained position measurements from Section IV-A, a Kalman
 342 Filter is implemented. The measurements from the flight controller are
 343 used in the prediction step.

Prediction Step:

$$\begin{aligned}\hat{x}_t &= F_t \cdot \hat{x}_{t-1} + G_t \cdot u_t \\ P_t &= F_t \cdot P_{t-1} \cdot F_t^\top + Q_t\end{aligned}\tag{1}$$

where:

$$\hat{x}_t = \begin{bmatrix} x \\ y \end{bmatrix}, \quad F_t = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix},$$

$$G_t = \begin{bmatrix} dt & 0 & \frac{dt^2}{2} & 0 & -dt & 0 \\ 0 & dt & 0 & \frac{dt^2}{2} & 0 & -dt \end{bmatrix}, \quad u_t = \begin{bmatrix} v_x \\ v_y \\ a_x \\ a_y \\ v_{x,\text{landing}} \\ v_{y,\text{landing}} \end{bmatrix}$$

$$P_0 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad Q_t = \begin{bmatrix} dt & 0 \\ 0 & dt \end{bmatrix} \cdot \sigma_{\text{IMU}}^2$$

Georgios: What are the value of the initial covariance matrix?

- x, y is the position of the UAV in the landing site's local coordinate system (aligned with NED frame),
- dt is the time interval the PX4's EKF outputs the v_x, v_y data, usually around 33ms,
- v_x, v_y is the linear velocity of the UAV in NED frame, estimated from the PX4's EKF,
- a_x, a_y is the linear acceleration of the UAV in NED frame, estimated on the companion computer from differentiating the velocities.
- $v_{x,\text{landing}}, v_{y,\text{landing}}$ is the linear velocity of the moving platform in NED frame, estimated from the velocity estimator for the moving platform,
- σ_{IMU}^2 is the variance of the velocities based estimated by the PX4's EKF.

360 In the correction step the observed measurements from the downward
 361 looking camera will be used to correct and update the predicted X, Y
 362 position of the UAV. However due to the computation time needed to
 363 detect the landing marker, that incoming measurement will be delayed by
 364 some time dt_{obs} . During that time dt_{obs} the UAV will be relocated by an
 365 interval (dx_{obs}, dy_{obs}) . To compensate that extra displacement, the interval
 366 (dx_{obs}, dy_{obs}) is calculated and added to the observed measurements.

Correction step:

$$\begin{aligned}
 e_t &= z_t - H_t \cdot \hat{x}_t \\
 S_t &= H_t \cdot P_t \cdot H_t^\top + R_t \\
 K_t &= P_t \cdot H_t^\top \cdot S_t^{-1} \\
 \hat{x}_t &= \hat{x}_t + K_t \cdot e_t \\
 P_t &= (I - K_t \cdot H_t) \cdot P_t
 \end{aligned} \tag{2}$$

where:

$$z_t = \begin{bmatrix} x_{obs} + dx_{obs} \\ y_{obs} + dy_{obs} \end{bmatrix}, \quad H_t = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad R_t = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} * \sigma_{obs}^2$$

$$dx_{obs} = vx_{mean_{obs}} \cdot dt_{obs}, \quad dy_{obs} = vy_{mean_{obs}} \cdot dt_{obs}$$

367 A velocity estimator is also constructed to determine the magnitude
 368 and direction of the moving platform's velocity vector. The magnitude is
 369 calculated by differentiating two sequential X, Y positions of the tractor,
 370 obtained by detecting the n-fold marker from the vertically facing camera
 371 as explained in section IV-A and taking into consideration the UAV's
 372 NED velocities according to the PX4's EKF. A low-pass filter is used to
 373 provide a smooth estimation of the velocity's magnitude by filtering out
 374 high frequency noise. High frequency noise can be caused by oscillations
 375 of the UAV along with a fast update rate in the landing marker detection
 376 algorithm.

377 In the agricultural use-case the moving platform is likely to change its
 378 direction up to 180 degrees. Furthermore, it is assumed that the moving
 379 platform is a nonholonomic system, like a tractor. To compensate for

Where do these position
measurements come
from? the tractor gps?

Fix loose sentence here.

sudden turns, the moving platform's yaw orientation will be taken into account. Based on the velocity's magnitude and moving platform's yaw orientation, the moving platform's velocity $(v_{x,\text{landing}}, v_{y,\text{landing}})$ in *NED* frame can be obtained.

C. Wind Disturbances

The exact and accurate estimation of the applied wind forces on a body is a complex matter studied by the field of fluid dynamics. We use a simplified approach, as follows. We assume that the wind will be applied on an area of 0.09 m^2 . Such an area is emulating a UAV with extra payload attached on its frame. Two different wind speeds of 8 m/s and 12 m/s will be used to calculate the applied wind forces on that area. The wind forces are considered to be applied on the center of gravity of the UAV with direction parallel to the ground.

The magnitude of the applied force, corresponding to a certain wind speed is calculated by the following equations [28], [29]:

$$\begin{aligned} F &= p_d \cdot A \\ p_d &= \frac{\rho \cdot v^2}{2} \end{aligned} \quad (3)$$

where F is the force, p_d is the dynamic pressure, A is the area of the applied pressure, ρ is the density of the air (around 1.2 kg/m^3), v is the wind velocity in m/s .

By solving the above equations the applied force on the UAV is found to be 3.45 N for an 8 m/s wind speed, and 7.76 N for a 12 m/s wind speed. The wind forces will be applied on the UAV in Gazebo simulation, to test the performance of the whole system. This will be done by constructing a program that will apply the wind forces on the virtual UAV.

To simulate a wind pattern the wind direction and magnitude must be defined. The wind direction is assumed to remain the same for the whole duration of the experiments. That direction vector is defined as $[0.8, 0.2]$ according to Gazebo's (x, y) axes. The magnitude of the wind is calculated as follows:

- 408 • An initial wind force of either 3.45 N or 7.76 N is chosen.
- 409 • An update cycle of 5.5 s is chosen between two different applied
- 410 forces.
- 411 • A random float number between 0.9 and 1.2 is chosen at every update
- 412 cycle. This number is defined as *Random_noise*.
- 413 • The applied wind force is calculated as:

$$414 \quad force_x = -x_{norm} * force * Random_noise$$

$$415 \quad force_y = -y_{norm} * force * Random_noise$$

416 where $x_{norm} = 0.8$, $y_{norm} = 0.2$

417 For an altitude of 6.0 m and above, $force = force_{init}$ where $force_{init}$

418 is either 3.45 N or 7.76 N

- 419 • For a UAV altitude of 6.0 m and below, the wind force decreases as:
- 420 $force = force_{init} * (altitude/6.0)$
- 421 • For a UAV altitude of 3 m and below, $force = 0.5N$.

422 This model is created to simulate different scaling in the wind's magnitude

423 and will be used for all simulated tests.

424 V. EVALUATION

425 We evaluate our approach in terms of the quality of the overall func-

426 tionality and the energy efficiency of the algorithms. All algorithms are

427 executed on an embedded companion computer interfaced to a simulation

428 running on a standard computer.

429 A. Use case: agricultural safety

430 We evaluate our approach based on a simulated use case where a multi-

431 rotor UAV identifies hazardous objects around a moving platform, and

432 lands on the moving platform to recharge. No communication link is

433 considered between the moving platform and the UAV and no GNSS

434 positioning is assumed to be available. The system can thus be considered

435 as a fallback for fault-tolerance.



Fig. 3. On the left, a top view of the Gazebo scene. On the right, a view of the UAV attempting a landing on the moving platform

Object detection and classification is performed by feeding the input image from the downward facing camera into a *YOLOv3-tiny* CNN [30] implemented in ROS [31]. Four different classes are selected: cars, humans, tractors, and cows. Based on the CNN's predictions and onboard sensors, the UAV maps the detected objects onto a 2D map.

A simulated field is created in Gazebo with objects placed in random positions and orientations as seen in Figure 3.

Pre-trained weights for *YOLOv3-tiny*, were initially used but the performance on detecting objects from a downward facing camera was not satisfactory. Since a dataset for detecting the above four classes from a top view was not available, we created an artificial dataset based on the Gazebo models. The dataset consists of 1200 images, 300 for each class. We trained the *YOLOv3-tiny* by its default training parameters for 5000 epochs, for an input image size of 416×416 pixels.

B. Experimental setup

All experiments are performed in Gazebo simulation under Ubuntu 18.04 and ROS Melodic on a i7-8550U 1.8GHz (4.0GHz Boost), 8GB DDR4 laptop. The *PX4* Software In The Loop (SITL) Firmware v1.10.2 is used as the flight controller and the *IRIS* quadcopter is used as the UAV platform. A vertically facing RGB camera is placed on the UAV providing a 640×480 pixel image at 10 fps. An *Nvidia Jetson Nano* with Ubuntu 18.04 and ROS Melodic is used as the UAV's companion computer. All the

computer vision, guidance, and control algorithms execute on the *Nvidia Jetson Nano*, similar to how they would be deployed if the Nano was a companion computer in a physical drone. Energy profiling is performed directly on the *Nvidia Jetson Nano* using `powprofiler` as outlined in Section III-B.

Two groups of experiments are conducted. The first group evaluates the energy consumption and QoS of the *tracking* mode and the second group evaluates the energy consumption and QoS of the *landing* mode. For both groups, the experiments start with the tractor moving at a constant speed of 0.3 m/s, according to a square path similar to that of a plowing tractor, and the UAV taking off and hovering at an altitude of 25 m. After reaching the desired altitude, the UAV starts searching for the landing marker in the image frame. Once the landing marker is detected, the UAV commences its actions.

In *tracking* mode, the UAV will follow the moving platform at a fixed altitude and use its vertically facing camera to map the environment, while in *landing* mode the UAV will follow the moving platform and gradually lower its altitude until it lands on it. Both *tracking* and *landing* modes are tested under three different cases of no wind disturbances, wind disturbances of 8 m/s and wind disturbances of 12 m/s according to the wind model described in Section IV-C.

C. Results

The first group of experiments was conducted to test the *tracking* mode and evaluate its energy efficiency and QoS. For the energy evaluation, eight tests were executed for different fps rates for the *YOLOv3-tiny* ROS node (4fps, 1fps, 0.5fps, 0.1fps) and the *landing marker* detection ROS node (10fps, 0.5fps) as seen in Figure 4. For a 4fps update rate for *YOLOv3-tiny*, a power consumption of 6.30 W is observed while for a 1fps and 0.1fps update rates, the power consumption drops to 4.8 W and 3.9 W accordingly. By reducing the update rate for *landing marker*




Fig. 4. Power consumption during tracking mode. YOLOv3 Tiny fps: 0=0.1fps, 1=0.5fps, 2=1fps, 3=4fps.

490 detection, from 10fps to 0.5fps, a further power saving of 0.15 W- 0.19 W
 491 is achieved.

492 For the QoS evaluation, twelve tests were executed for different fps
 493 rates for the *YOLOv3-tiny* ROS node (4fps, 0.1fps) and for the *landing*
 494 *marker* detection ROS node (10fps, 0.5fps) under three different cases
 495 of wind disturbances. The QoS is determined as the number of correctly
 496 detected objects. An object is considered to be correctly detected if it is
 497 detected within a distance of 2 m from the its actual position and classified
 498 with the correct class. The detection results can be seen in Figure 5. The
 499 best results were obtained for a high fps update rate in *YOLOv3-tiny* and
 500 *landing marker* detection, under no wind disturbances, since 28 out of the
 501 32 objects were detected. Nevertheless, for the same high fps values but
 502 with a wind speed of 12 m/s, only 18 out of 32 objects were correctly
 503 detected.

504 The second group of experiments was conducted to test the *landing*
 505 mode and evaluate its energy efficiency and QoS. For the energy evalua-
 506 tion, eight tests were executed for different fps rates for the *landing marker*
 507 detection ROS node (10fps, 2fps, 1fps, 0.5fps) using two different cases.
 508 In the first case, the kernel size remains fixed at 22×22 pixels and in the
 509 second case an adaptive selection kernel size algorithm is used. The landing




Fig. 5. Number of correctly detected objects under different conditions. Grey color denotes a 0.1fps and black color denotes a 4fps update rate in *YOLOv3-tiny* ROS node. Circles denote a 0.5fps and triangles denote a 10fps update rate in the *landing marker* detection ROS node.

time is also taken under consideration as seen in Figure 6. The largest difference in the power consumption is 0.14 W and is observed between the 0.5fps and the 10fps update rate for the *landing marker* detection. However the landing time is reduced by 30 s when using the 10fps update rate. The adaptive kernel size in most cases outperforms the fixed kernel size by 0.11 W. Furthermore the adaptive kernel size can compensate for marker occlusions which will increase the overall robustness of the system.

For the QoS evaluation, twelve tests were executed for different fps rates for the *landing marker* detection ROS node (10fps, 2fps, 1fps, 0.5fps) under three different cases of wind disturbances. The QoS is determined as the mean squared error (MSE) between the predicted position of the moving platform, from the navigation block, and the moving platform's actual position. Four different altitude bins were used as seen in Figure 7. A large MSE of around 3 m^2 is observed for an altitude greater than 20 m for a 0.5fps rate, while an MSE close to zero is observed for an altitude of less than 5 m for an update rate of 10fps. Furthermore, wind disturbances do not seem to have an influence on the MSE. We believe the larger MSE for the y coordinates compared to the x coordinates is caused by a sudden change of the moving platform's direction on the y axis.

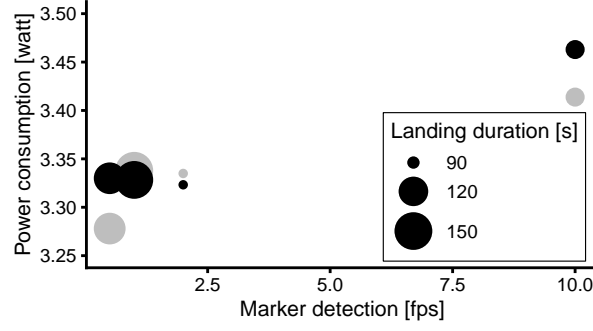


Fig. 6. Power consumption during landing mode. The black circles denote a fixed kernel size while the grey circles denote an adaptive kernel size.




Fig. 7. Position error during landing and how it depends on the marker detection rate at different altitudes and wind disturbances. Circles denote errors in the x direction and triangles errors in the y direction.

529 D. Discussion

530 The experiments show that both tracking mode and landing mode are
 531 supported by the system, in a simulated environment with a moving plat-
 532 form and random wind conditions. Moreover, the performance of both
 533 modes is sensitive to the QoS, with a high success rate of both modes at
 534 high QoS levels, and significantly lower performance at lower QoS levels.

535 The potential energy savings from having an energy-sensitive algorithm
 536 that can adapt its QoS by changing the fps values for the *YOLOv3-tiny* and

537 *landing marker* ROS nodes should be seen in relation to the total energy
 538 consumption of the UAV. As a concrete example, consider a DJI Phantom
 539 4 multirotor and a Sky-Watch Cumulus fixed-wing (the fixed-wing would
 540 need to circle while tracking and would need VTOL capabilities to land,
 541 but we nevertheless include it for comparison). We estimate¹ that the
 542 Phantom uses roughly 140 W while cruising whereas the Cumulus uses
 543 roughly 40 W while cruising. The maximal saving gained from changing
 544 the *YOLOv3-tiny* rate is $6.30W - 3.9W = 2.4W$ whereas the maximal
 545 saving gained from changing the *landing marker* rate is 0.2 W. For the
 546 Cumulus, there is thus a 6.5% potential energy saving, whereas the po-
 547 tential energy saving only is 1.9% for the Phantom. For the Cumulus this
 548 saving is considered large enough to significantly impact the flying time
 549 of the drone, with a total energy saving of 23.4 kJ. For the Cumulus, the
 550 potential saving from adapting the *landing marker* QoS is however only
 551 0.5%.

552 For the tracking mode, changing the *landing marker* rate provided a
 553 minor saving of 0.14 W, but at the cost of increased landing time. There-
 554 fore, although the higher-QoS computer vision algorithm is marginally
 555 more expensive by 0.14 W, the UAV will overall save energy because of
 556 a reduced flight time.

557 VI. CONCLUSION AND FUTURE WORK

558 In this paper we presented a robust, energy-sensitive, vision-based al-
 559 gorithm for autonomous tracking and landing in varying environmental
 560 conditions, by experimentally executing all the necessary algorithms on
 561 the *Nvidia Jetson Nano* companion computer. Our experiments show that
 562 the proposed computer vision algorithms for detecting the moving plat-
 563 form can be run at the highest QoS level with only a marginal energy
 564 overhead, whereas adapting the QoS level of *YOLOv3-tiny* CNN results in

¹From information on the respective product pages regarding battery capacity and maximal flight time.

a considerable power saving for the system as a whole. This power saving is significant if the system was executing on a fixed-wing UAV, but only marginal if executing on a multirotor UAV.

In terms of future work, we are interested in automatically adapting the QoS level to the available battery, and in testing this approach on a physical drone.

ACKNOWLEDGMENT

This work is supported and partly funded by the European Unions Horizon2020 research and innovation program under grant agreement No. 779882 (TeamPlay).

REFERENCES

- [1] F. G. Costa, J. Ueyama, T. Braun, G. Pessin, F. S. Osório, and P. A. Vargas, "The use of unmanned aerial vehicles and wireless sensor network in agricultural applications," in *2012 IEEE International Geoscience and Remote Sensing Symposium*. IEEE, 2012, pp. 5045–5048.
- [2] E. Salami, C. Barrado, and E. Pastor, "Uav flight experiments applied to the remote sensing of vegetated areas," *Remote Sensing*, vol. 6, no. 11, pp. 11 051–11 081, 2014.
- [3] A. Seewald, H. Garcia de Marina, H. S. Midtby, and U. P. Schultz, "Mechanical and computational energy estimation of a fixed-wing drone," in *Proceedings of the 2020 Fourth IEEE International Conference on Robotic Computing (IRC)*. IEEE, 2020, p. to appear.
- [4] S. Saripalli, J. F. Montgomery, and G. S. Sukhatme, "Vision-based autonomous landing of an unmanned aerial vehicle," in *Proceedings 2002 IEEE International Conference on Robotics and Automation (Cat. No. 02CH37292)*, vol. 3. IEEE, 2002, pp. 2799–2804.
- [5] H. Yuan, C. Xiao, S. Xiu, W. Zhan, Z. Ye, F. Zhang, C. Zhou, Y. Wen, and Q. Li, "A hierarchical vision-based localization of rotor unmanned aerial vehicles for autonomous landing," *International Journal of Distributed Sensor Networks*, vol. 14, no. 9, 2018.
- [6] S. Saripalli and G. S. Sukhatme, "Landing on a moving target using an autonomous helicopter," in *Field and service robotics*. Springer, 2003, pp. 277–286.
- [7] D. Lee, T. Ryan, and H. J. Kim, "Autonomous landing of a vtol uav on a moving platform using image-based visual servoing," in *2012 IEEE international conference on robotics and automation*. IEEE, 2012, pp. 971–976.
- [8] O. Araar, N. Aouf, and I. Vitanov, "Vision based autonomous landing of multirotor uav on moving platform," *Journal of Intelligent & Robotic Systems*, vol. 85, no. 2, pp. 369–384, 2017.

- 600 [9] E. Olson, "Apriltag: A robust and flexible visual fiducial system," in *2011 IEEE*
601 *International Conference on Robotics and Automation*. IEEE, 2011, pp. 3400–3407.
- 602 [10] X. Chen, S. K. Phang, M. Shan, and B. M. Chen, "System integration of a vision-guided
603 uav for autonomous landing on moving platform," in *2016 12th IEEE International*
604 *Conference on Control and Automation (ICCA)*. IEEE, 2016, pp. 761–766.
- 605 [11] H. Lee, S. Jung, and D. H. Shim, "Vision-based uav landing on the moving vehicle," in
606 *2016 International conference on unmanned aircraft systems (ICUAS)*. IEEE, 2016,
607 pp. 1–7.
- 608 [12] T. Yang, Q. Ren, F. Zhang, B. Xie, H. Ren, J. Li, and Y. Zhang, "Hybrid camera array-
609 based uav auto-landing on moving ugv in gps-denied environment," *Remote Sensing*,
610 vol. 10, no. 11, p. 1829, 2018.
- 611 [13] P. H. Nguyen, M. Arsalan, J. H. Koo, R. A. Naqvi, N. Q. Truong, and K. R. Park,
612 "Lightdenseyolo: A fast and accurate marker tracker for autonomous uav landing by
613 visible light camera sensor on drone," *Sensors*, vol. 18, no. 6, p. 1703, 2018.
- 614 [14] S. Kyristis, A. Antonopoulos, T. Chanielakis, E. Stefanakis, C. Linardos, A. Tripolitsi-
615 otis, and P. Partsinevelos, "Towards autonomous modular uav missions: The detection,
616 geo-location and landing paradigm," *Sensors*, vol. 16, no. 11, p. 1844, 2016.
- 617 [15] L. Nardi, B. Bodin, M. Z. Zia, J. Mawer, A. Nisbet, P. H. Kelly, A. J. Davison,
618 M. Luján, M. F. O'Boyle, G. Riley *et al.*, "Introducing slambench, a performance and
619 accuracy benchmarking methodology for slam," in *2015 IEEE International Conference*
620 *on Robotics and Automation (ICRA)*. IEEE, 2015, pp. 5783–5790.
- 621 [16] A. Seewald, U. P. Schultz, E. Ebeid, and H. S. Midtiby, "Coarse-grained computation-
622 oriented energy modeling for heterogeneous parallel embedded systems," *International*
623 *Journal of Parallel Programming*, pp. 1–22, 2019.
- 624 [17] A. Sadrpour, J. Jin, and A. G. Ulsoy, "Experimental validation of mission energy pre-
625 diction model for unmanned ground vehicles," in *2013 American Control Conference*.
626 IEEE, 2013, pp. 5960–5965.
- 627 [18] —, "Mission energy prediction for unmanned ground vehicles using real-time
628 measurements and prior knowledge," *Journal of Field Robotics*, vol. 30, no. 3, pp.
629 399–414, 2013.
- 630 [19] J. Morales, J. L. Martinez, A. Mandow, A. J. García-Cerezo, and S. Pedraza, "Power
631 consumption modeling of skid-steer tracked mobile robots on rigid terrain," *IEEE*
632 *Transactions on Robotics*, vol. 25, no. 5, pp. 1098–1108, 2009.
- 633 [20] Y. Mei, Y.-H. Lu, Y. C. Hu, and C. G. Lee, "Deployment of mobile robots with energy
634 and timing constraints," *IEEE Transactions on robotics*, vol. 22, no. 3, pp. 507–522,
635 2006.
- 636 [21] —, "A case study of mobile robot's energy consumption and conservation tech-
637 niques," in *ICAR'05. Proceedings., 12th International Conference on Advanced*
638 *Robotics, 2005*. IEEE, 2005, pp. 492–497.
- 639 [22] —, "Energy-efficient motion planning for mobile robots," in *IEEE International*

- 640 *Conference on Robotics and Automation, 2004. Proceedings. ICRA'04. 2004*, vol. 5.
 641 IEEE, 2004, pp. 4344–4349.
- 642 [23] V. Berenz, F. Tanaka, and K. Suzuki, “Autonomous battery management for mobile
 643 robots based on risk and gain assessment,” *Artificial Intelligence Review*, vol. 37, no. 3,
 644 pp. 217–237, 2012.
- 645 [24] Y. Feng, C. Zhang, S. Baek, S. Rawashdeh, and A. Mohammadi, “Autonomous landing
 646 of a uav on a moving platform using model predictive control,” *Drones*, vol. 2, no. 4,
 647 p. 34, 2018.
- 648 [25] D. Falanga, A. Zanchettin, A. Simovic, J. Delmerico, and D. Scaramuzza, “Vision-
 649 based autonomous quadrotor landing on a moving platform,” in *2017 IEEE International
 650 Symposium on Safety, Security and Rescue Robotics (SSRR)*. IEEE, 2017, pp. 200–207.
- 651 [26] Henrik Skov Midthiby, “N-fold marker tracker repository,” [https://github.com/
 652 henrikmidthiby/MarkerLocator](https://github.com/henrikmidthiby/MarkerLocator), 2015.
- 653 [27] OpenCV, “Detection of ArUco markers,” [https://docs.opencv.org/3.4/d5/dae/tutorial_
 654 aruco_detection.html](https://docs.opencv.org/3.4/d5/dae/tutorial_aruco_detection.html), 2020.
- 655 [28] NASA, “Dynamic pressure (NASA),” [https://www.grc.nasa.gov/WWW/K-12/airplane/
 656 dynpress.html](https://www.grc.nasa.gov/WWW/K-12/airplane/dynpress.html), 2020.
- 657 [29] J. D. Anderson Jr, *Fundamentals of aerodynamics*. Tata McGraw-Hill Education, 2010.
- 658 [30] J. Redmon and A. Farhadi, “Yolov3: An incremental improvement,” *arXiv*, 2018.
- 659 [31] M. Bjelonic, “YOLO ROS: Real-time object detection for ROS,” [https://github.com/
 660 leggedrobotics/darknet_ros](https://github.com/leggedrobotics/darknet_ros), 2016–2018.